

Running game servers on Kubernetes

Mo Firouz

About me

- My name is Mo(hammad) Firouz!
- Originally Iranian - can speak Farsi with a funny accent.
- Software architect
- Co-founder of Heroic Labs

Heroic Labs

- Nakama, a realtime and social server for games.
- Completely open-source (and free!):
<https://github.com/heroiclabs/nakama>
- Realtime, server-authoritative multiplayer
- Flexible matchmaking system
- Extendable using Lua scripting
- Friends, Clans, Chat + more



Nakama

Things to discuss

1. Intro to Kubernetes
2. System overview
3. Two examples of game servers
 - Nakama server
 - Unity Headless servers

Intro to Kubernetes

*System for automating deployment, scaling, and management of **containerized** applications*

- Hardware provisioning - replaces Puppet, Chef and Ansible.
- Supervisor-like responsibilities - healthcheck and uptime
- Log aggregation + metrics + scheduled jobs + (much) more

Intro to Kubernetes

Some rules to follow:

1. Log to stdout and stderr
2. Return non-zero for abnormal exits
3. Add healthcheck endpoint - ideally to /
4. Add readiness endpoint - ideally to same as healthcheck
5. Use cmd args and env variables, not config files

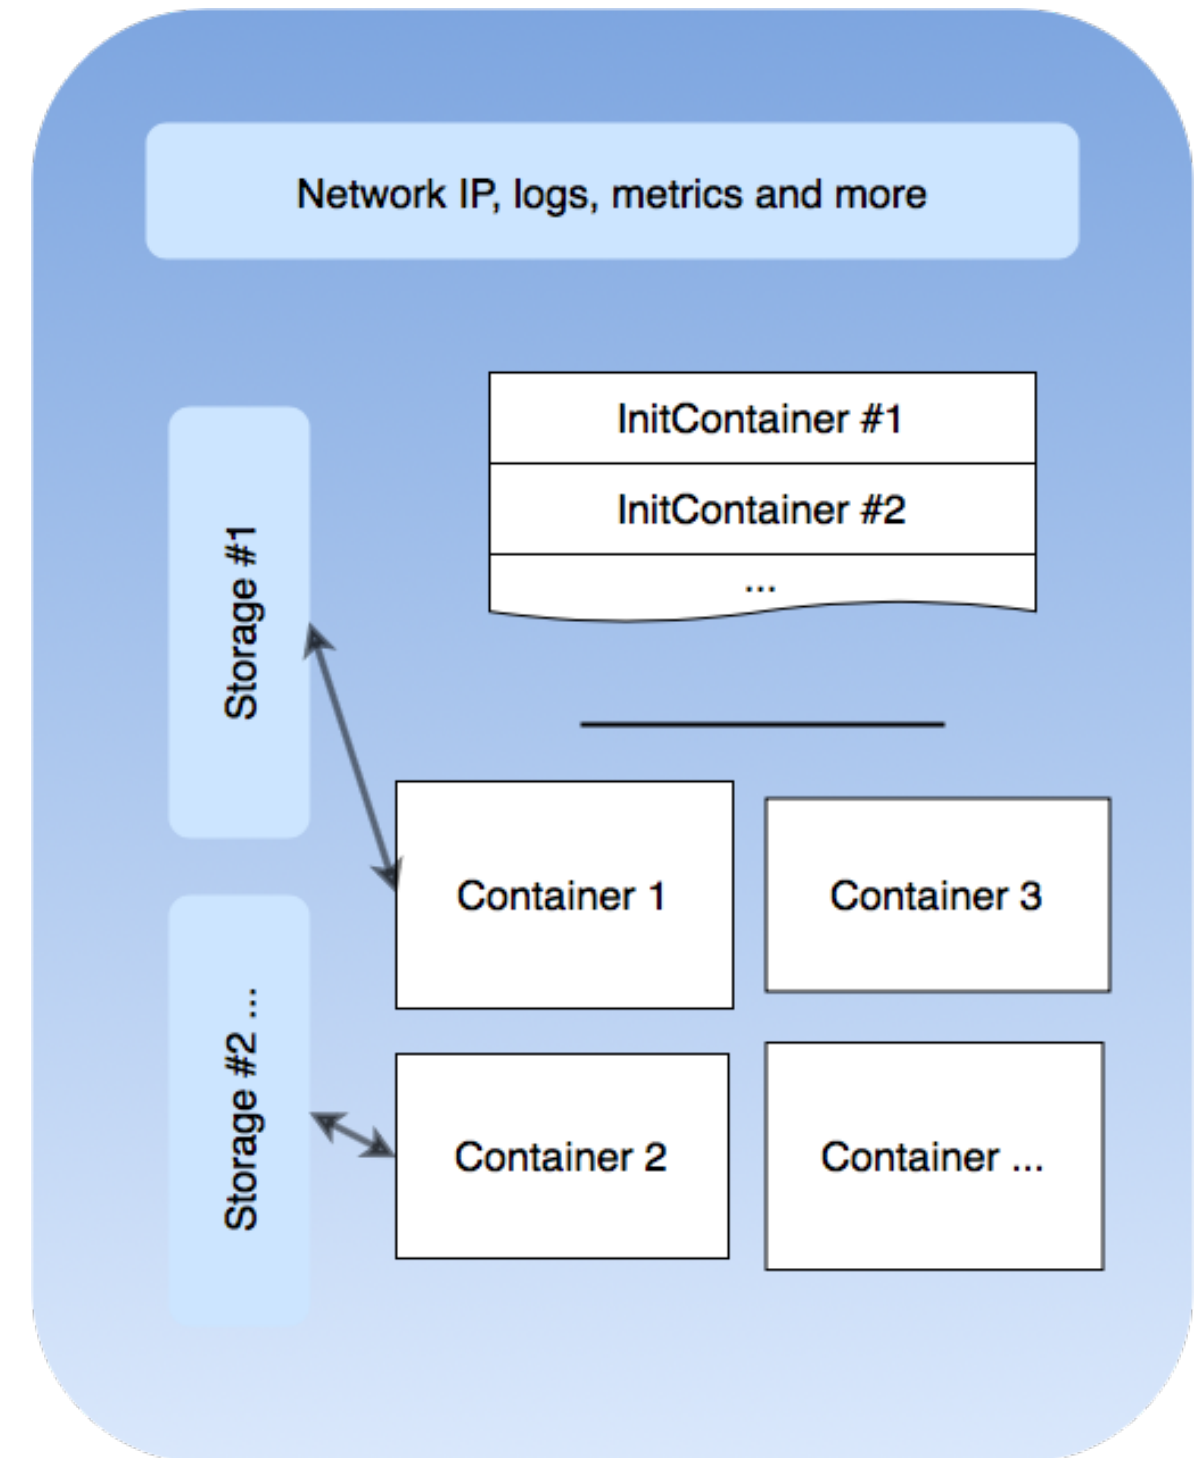
Intro to Kubernetes

Kubernetes components

1. Pod
2. Controllers
3. Service + Ingress
4. Nodes
5. Autoscaling

Pod

- Wrap one or more containers
- Pods are ephemeral
 - can be dynamically created / deleted
- Unique network cluster-IP (but not permanent)
- Storage resource(s)



Pod

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    app: nakama
spec:
  volumes:
  - name: nakama-volume
    emptyDir:
containers:
- name: nakama
  image: heroiclabs/nakama
  command: "/bin/sh exec nakama"
  ports:
  - containerPort: 7350
    name: api
  volumeMounts:
  - mountPath: /nakama-data
    name: nakama-volume
```

Service

- A way to access pods
- Services have DNS names
 - `controller.namespace.service.cluster.local`
 - `controller.namespace`
- Expose applications internally and externally

Service

- ClusterIP: cluster-internal IP
- NodePort: Expose port (mapping) on the physical node
- LoadBalancer: Cloud provider's load balancer
- ExternalName: CNAME without any proxying

Service

```
apiVersion: v1
kind: Service
metadata:
  name: nakamaservice
  labels:
    service: nakama
spec:
  selector:
    app: nakama
  type: NodePort
  ports:
    - port: 80
      targetPort: 7350
      name: api
```

Ingress

- Very similar to LoadBalancer service
- Only allowing HTTP(S) traffic into the cluster
- SSL Termination
- Usually configures Cloud Provider traffic

Ingress

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: nakamaingress
  labels:
    ingress: nakama
  annotations:
    kubernetes.io/ingress.allow-http: "false"
spec:
  tls:
    - secretName: ingress-cert-secret
  backend:
    serviceName: nakama
    servicePort: 80
```

Controllers

- Replicating homogeneous set of pods
- Ensures healthy pods are running
- Deployment:
 - Manage infrastructure state from current to desired
 - Rolling update
 - Scale up/down pod replicas
 - History and rollback

Controller

- StatefulSet:
 - Very similar to Deployment
 - Sticky pod IDs - sequential guaranteed ordering
 - Disk and Network resources will stay the same across roll outs
- DaemonSets, (Cron)Jobs, and more...

Controller

```
apiVersion: apps/v1beta2
kind: StatefulSet
metadata:
  name: nakamastatefulset
spec:
  serviceName: "nakamaservice" # This must exist before Statefulset is created
  revisionHistoryLimit: 1
  replicas: 2
  podManagementPolicy: OrderedReady # or Parallel
  updateStrategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: nakama # this must matches below
  template:
    metadata:
      labels:
        app: nakama
    # ...rest of pod template...
```

Nodes

- Assign special pods to special nodes

```
nodeSelector:  
  disktype: ssd
```

- Use `affinity` and `anti-affinity` to indicate soft/hard requirement for node attractiveness
- Use `taint` and `tolerations` to repel pods from nodes
- Use both to schedule pods only on nodes with special hardware
- Ensure to use `resource requests` and `limits` for pods to limit noisy neighbour problems.

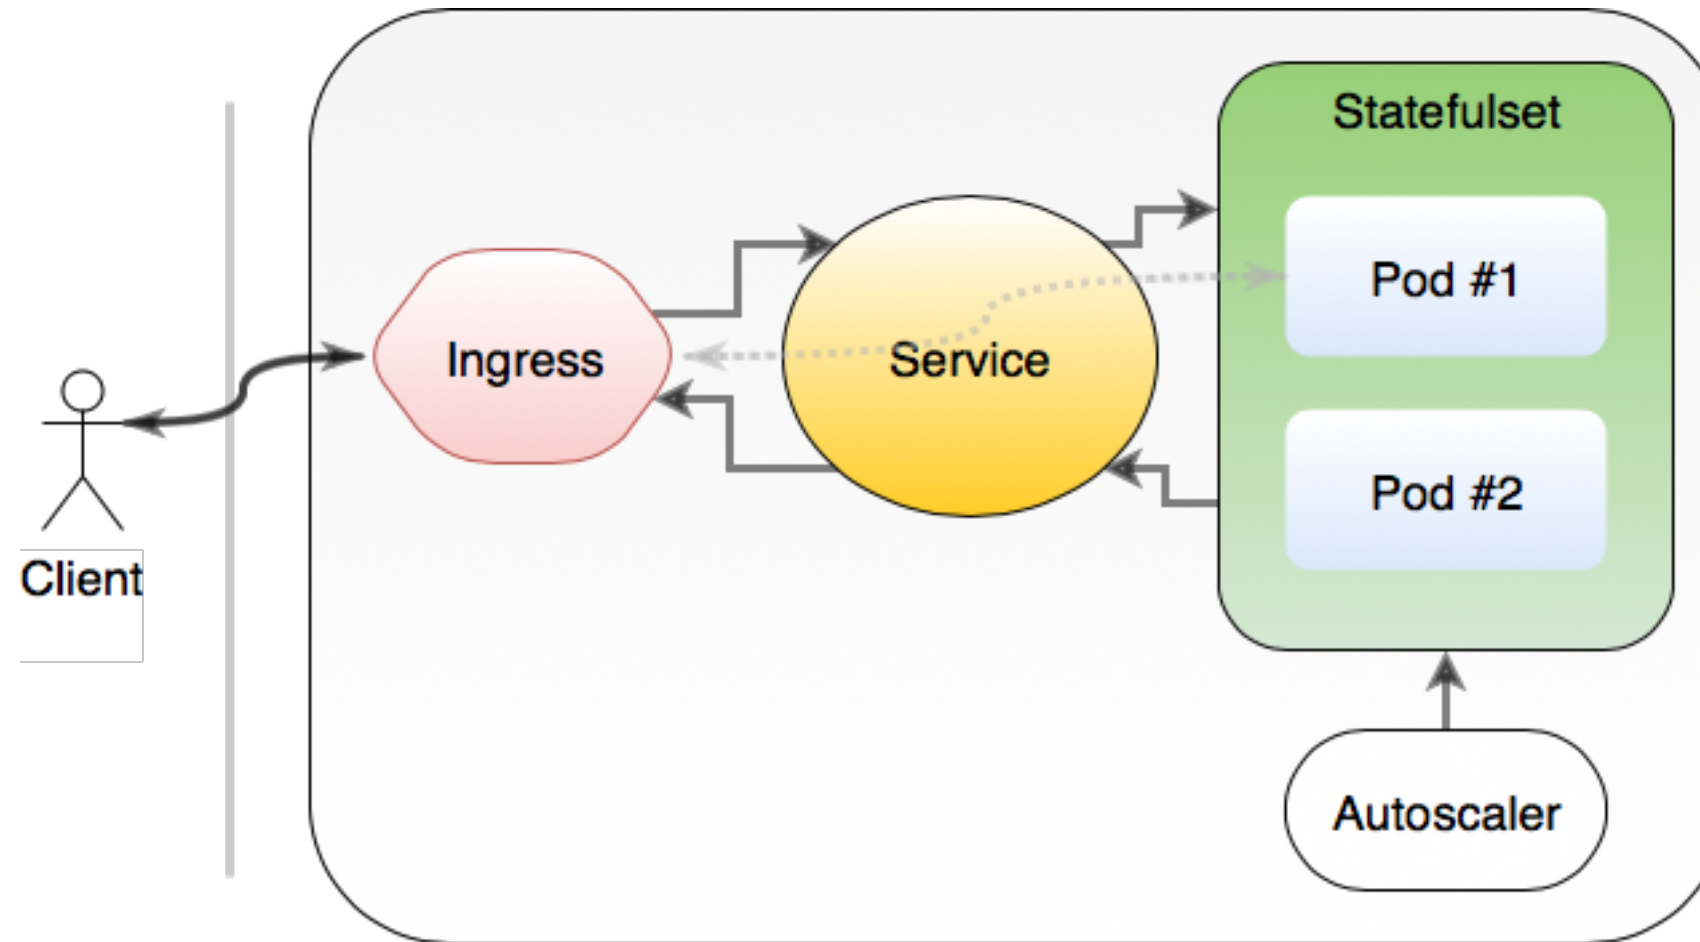
Nodes

```
affinity:
  podAntiAffinity:
    preferredDuringSchedulingIgnoredDuringExecution:
      - weight: 100
        podAffinityTerm:
          topologyKey: kubernetes.io/hostname
          labelSelector:
            matchExpressions:
              - key: app
                operator: In
                values:
                  - nakama
```

Autoscaling

```
apiVersion: autoscaling/v2beta1
kind: HorizontalPodAutoscaler
metadata:
  name: nakamascaler
  labels:
    autoscaler: nakama
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: StatefulSet
    name: nakamastatefulset
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      targetAverageUtilization: 50
      name: cpu
```

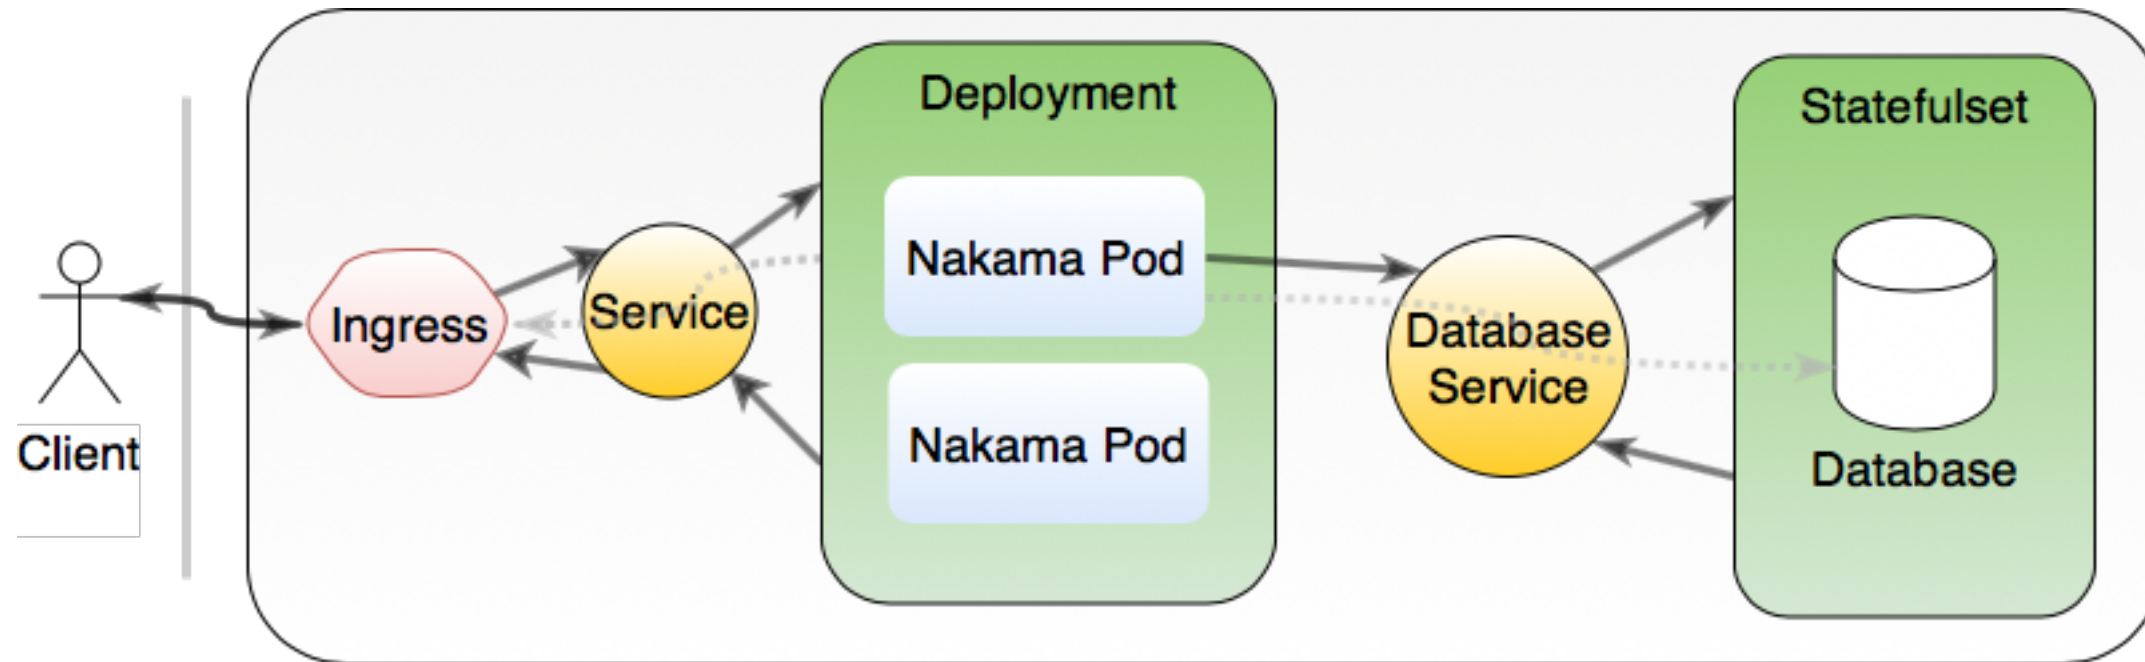
System overview



Nakama game server




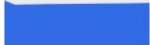




- Realtime multiplayer
- Server-authoritative multiplayer
- Matchmaking
- Custom code runtime
- Presence system
- Friends, Groups/clans
- Realtime chat, notifications
- User accounts, authentication

Nakama game server



Nakama game server

Pods





Name ▾	Node	Status ▾	Restarts	Age ▾	CPU (cores)	Memory (bytes)		
✓ nakama-0	nakama-n1s1-7a45d00d-b0n3	Running	0	10 hours	 0.005	 64.531 Mi	≡	⋮
✓ cockroachdb-1	db-n1s2-a7f0d49a-f90p	Running	0	16 days	 0.019	 1.016 Gi	≡	⋮
✓ cockroachdb-0	db-n1s2-a7f0d49a-8cb2	Running	0	16 days	 0.016	 973.879 Mi	≡	⋮
✓ cockroachdb-2	db-n1s2-a7f0d49a-4ts8	Running	0	16 days	 0.017	 1,013.605 Mi	≡	⋮

Stateful Sets

Name ▾	Labels	Pods	Age ▾	Images		
✓ nakama	statefulset: nakama	1 / 1	a month	heroiclabs/nakama:2.0.1	≡	⋮
✓ cockroachdb	statefulset: cockroachdb	3 / 3	2 months	cockroachdb/cockroach:v2.0.1	≡	⋮

Nakama game server

Ingresses						
Name	Endpoints			Age		
nakama	35.186.228.0			2 months		

Services						
Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age	
 nakama	service: nakama	10.23.250.87	nakama.tgc:7350 TCP nakama.tgc:30570 TCP	-	a month	
 cockroachdb	service: cockroachdb	None	cockroachdb.tgc:26257 T cockroachdb.tgc:0 TCP cockroachdb.tgc:8080 TC cockroachdb.tgc:0 TCP	-	2 months	

Unity headless server

```
/Applications/Unity/Unity.app/Contents/MacOS/Unity -quit -batchmode -nographics \  
-logFile .unity-build.log \  
-projectPath $(pwd)/unity-fastpacedmultiplayer \  
-executeMethod BuildTools.QuickBuildLinux
```

```
public static void QuickBuildLinux()  
{  
    BuildPlayerOptions opts = new BuildPlayerOptions  
    {  
        options = BuildOptions.Development & BuildOptions.EnableHeadlessMode,  
        locationPathName = "/Users/mo/Desktop/server",  
        target = BuildTarget.StandaloneLinux64  
    };  
  
    BuildPipeline.BuildPlayer(opts);  
}
```

Unity headless server

```
metadata:
  generateName: "unity-"
  labels:
    app: unity-server
spec:
  hostNetwork: true
  containers:
    - image: mofirouz/unity-fastpacedmultiplayer:0.0.1
      name: unity-server
      imagePullPolicy: Always
      command:
        - "/bin/sh"
        - "-ecx"
        - >
          exec ./server -logFile /dev/stdout
  ports:
    - containerPort: 7777
      name: server-api
```

Feel free to send questions;

email: mo@heroiclabs.com

twitter: [@mofirouz](https://twitter.com/mofirouz)

skype: mo.firouz

Talk available on GitHub:

<https://github.com/mofirouz/gameservers-kubernetes>